

Case Study: JVM



Virtual Machines

What is a machine?


- does something (...useful)
- programmable
- concrete (hardware)

What is a virtual machine?

- a machine that is not concrete
- a software emulation of a physical computing environment

Reality is somewhat fuzzy!

Is a Pentium-II a machine?



Hardware and software are
logically equivalent
(A. Tanenbaum)

Virtual Machine, Intermediate Language

- Pascal P-Code (1975)
 - stack-based processor
 - strong type machine language
 - compiler: one front end, many back ends
 - UCSD Apple][implementation, PDP 11, Z80
- Modula M-Code (1980)
 - high code density
 - Lilith as microprogrammed virtual processor
- JVM – Java Virtual Machine (1995)
 - Write Once – Run Everywhere
 - interpreters, JIT compilers, Hot Spot Compiler
- Microsoft .NET (2000)
 - language interoperability

JVM Case Study

- compiler (Java to bytecode)
- interpreter, ahead-of-time compiler, JIT
- dynamic loading and linking
- exception Handling
- memory management, garbage collection
- OO model with single inheritance and interfaces
- system classes to provide OS-like implementation
 - compiler
 - class loader
 - runtime
 - system

JVM: Type System

- Primitive types
 - byte
 - short
 - int
 - long
 - float
 - double
 - char

 - reference
 - boolean mapped to int
- Object types
 - classes
 - interfaces
 - arrays
- Single class inheritance
- Multiple interface implementation
- Arrays
 - anonymous types
 - subclasses of `java.lang.Object`

JVM: Java Byte-Code

Memory access

- tload / tstore
- ttload / ttstore
- tconst
- getfield / putfield
- getstatic / putstatic

Operations

- tadd / tsub / tmul / tdiv
- tshifts

Conversions

- f2i / i2f / i2l /
- dup / dup2 / dup_x1 / ...

Control

- ifeq / ifne / iflt /
- if_icmpeq / if_acmpeq
- invokestatic
- invokevirtual
- invokeinterface
- athrow
- treturn

Allocation

- new / newarray

Casting


- checkcast / instanceof

JVM: Java Byte-Code Example

bipush

Operation Push **byte**

Format



Forms *bipush* = 16 (0x10)

Operand Stack ... => ..., *value*

Description The immediate *byte* is sign-extended to an int *value*. That *value* is pushed onto the operand stack.

JVM: Machine Organization

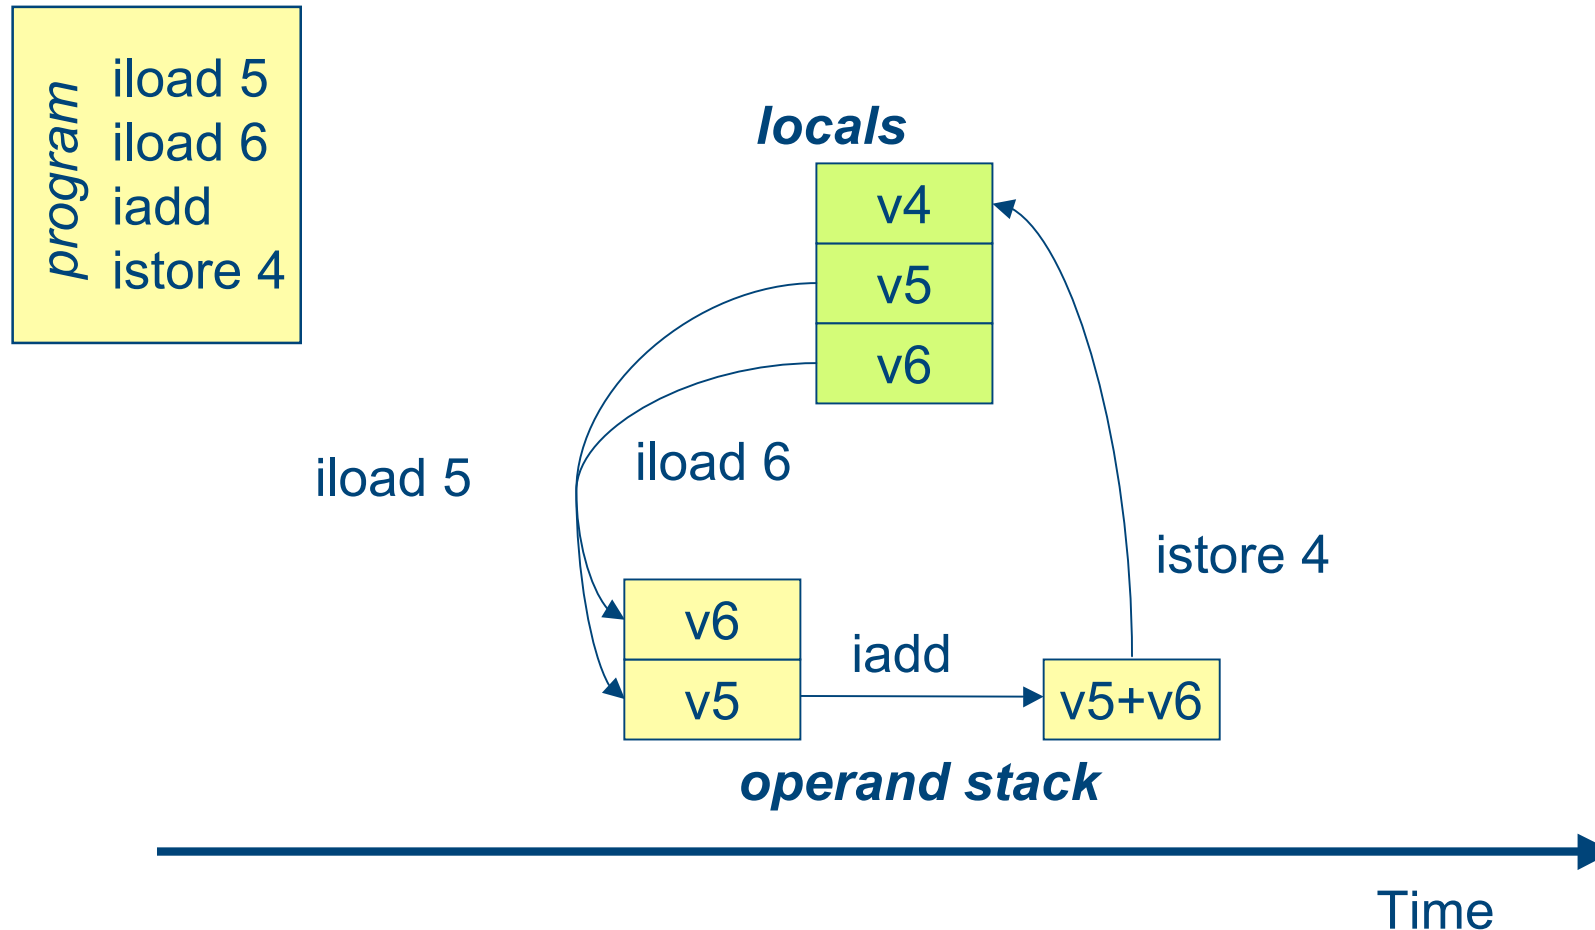
Virtual Processor

- stack machine
- no registers
- typed instructions
- no memory addresses, only symbolic names

Runtime Data Areas

- pc register
- stack
 - locals
 - parameters
 - return values
- heap
- method area
 - code
- runtime constant pool
- native method stack

JVM: Execution Example



JVM: Reflection

Load and manipulate *unknown* classes at runtime.

- java.lang.Class
 - getFields
 - getMethods
 - getConstructors
- java.lang.reflect.Field
 - setObject getObject
 - setInt getInt
 - setFloat getFloat
 -
- java.lang.reflect.Method
 - getModifiers
 - invoke
- java.lang.reflect.Constructor

JVM: Reflection – Example

```
import java.lang.reflect.*;

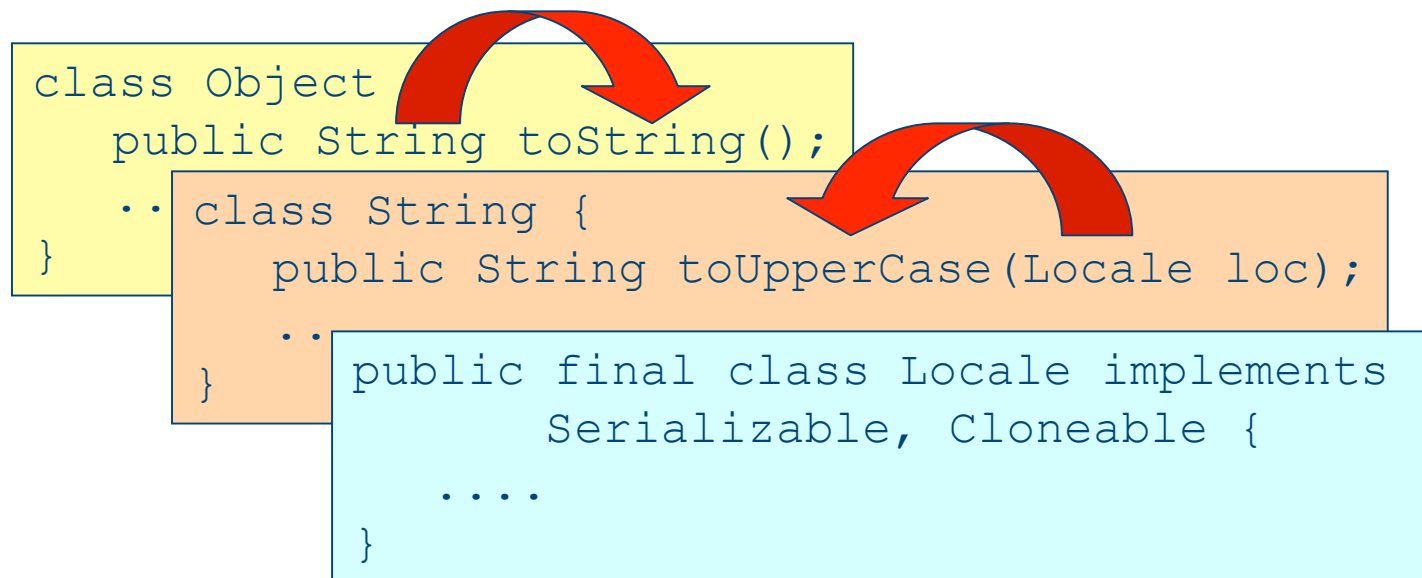
public class ReflectionExample {

    public static void main(String args[]) {
        try {
            Class c = Class.forName(args[0]);
            Method m[] = c.getDeclaredMethods();
            for (int i = 0; i < m.length; i++) {
                System.out.println(m[i].toString());
            }
        } catch (Throwable e) {
            System.err.println(e);
        }
    }
}
```

JVM: Java Weaknesses

Transitive closure of java.lang.Object contains

- 1.1 47
- 1.2 178
- 1.3 180
- 1.4 248
- 5 (1.5) 280
- classpath 0.03 299



JVM: Java Weaknesses

Class static initialization

- T is a class and an instance of T is created

```
T tmp = new T();
```

- T is a class and a static method of T is invoked


```
T.staticMethod();
```

- A nonconstant static field of T is used or assigned (field is not static, not final, and not initialized with compile-time constant)

```
T.someField = 42;
```

Problem

- circular dependencies in static initialization code



JVM: Java Weaknesses

```
interface Example {  
    final static String labels[] = {"A", "B", "C"}  
}
```

hidden static initializer:

```
labels = new String[3];  
labels[0] = "A"; labels[1] = "B"; labels[2] = "C";
```

Warning:

- in Java `final` means write-once!
- interfaces may contain code

JVM: Memory Model

- The JVM specs define a **memory model**:
 - defines the relationship between variables and the underlying memory
 - meant to guarantee the same behavior on every JVM
- The compiler is allowed to **reorder** operation unless **synchronized** or **volatile** is specified.

JVM: Reordering

- read and writes to ordinary variables can be reordered.

```
public class Reordering {  
    int x = 0, y = 0;  
  
    public void writer() {  
        x = 1;  
        y = 2;  
    }  
  
    public void reader() {  
        int r1 = y;  
        int r2 = x;  
    }  
}
```


JVM: Memory Model

- synchronized: in addition to specify a monitor it defines a **memory barrier**:
 - acquiring the lock implies an invalidation of the caches
 - releasing the lock implies a write back of the caches
- synchronized blocks on the same object are **ordered**.
- order among accesses to volatile variables is guaranteed (but **not** among volatile and other variables).

JVM: Double Checked Lock

Singleton

```
public class SomeClass {  
  
    private Resource resource = null;  
  
    public Resource synchronized getResource() {  
  
        if (resource == null) {  
            resource = new Resource();  
        }  
  
        return resource;  
    }  
}
```

JVM: Double Checked Lock

Double checked locking

```
public class SomeClass {  
  
    private Resource resource = null;  
  
    public Resource                getResource() {  
        if (resource == null) {  
            synchronized {  
                if (resource == null) {  
                    resource = new Resource();  
                }  
            }  
        }  
        return resource;  
    }  
}
```

JVM: Double Checked Lock

Thread 1

```
public class SomeClass {  
  
    private Resource resource  
        = null;  
  
    public Resource getResource() {  
        if (resource == null) {  
            synchronized {  
                if (resource == null) {  
                    resource =  
                        new Resource();  
                }  
            }  
        }  
        return resource;  
    }  
}
```

Thread 2

```
public class SomeClass {  
  
    private Resource resource  
        = null;  
  
    public Resource getResource() {  
        if (resource == null) {  
            synchronized {  
                if (resource == null) {  
                    resource =  
                        new Resource();  
                }  
            }  
        }  
        return resource;  
    }  
}
```

The object is
instantiated
but not yet initialized!

JVM: Immutable Objects are not Immutable

- Immutable objects:
 - all types are primitives or references to immutable objects
 - all fields are final
- Example (simplified): `java.lang.String`
 - contains
 - an array of characters
 - the length
 - an offset
 - example: `s = "abcd"`, `length = 2`, `offset = 2`, `string = "cd"`

```
String s1 = "/usr/tmp"
```

```
String s2 = s1.substring(4); //should contain "/tmp"
```

- Sequence: `s2` is instantiated, the fields are initialized (to 0), the array is copied, the fields are written by the constructor.
- What happens if instructions are reordered?


JVM: Reordering Volatile and Nonvolatile Stores

- volatile reads and writes are totally ordered among threads
- but not among normal variables
- example

```
volatile boolean initialized = false;  
SomeObject o = null;
```

Thread 1

```
o = new SomeObject;  
initialized = true;
```



Thread 2

```
while (!initialized) {  
    sleep();  
}  
o.field = 42;
```

JVM: JSR 133

- Java Community Process
- Java memory model revision
- Final means final
- Volatile fields cannot be reordered

Java JVM: Execution

- Interpreted (e.g., Sun JVM)
 - bytecode instructions are interpreted sequentially
 - the VM emulates the Java Virtual Machine
 - slower
 - quick startup
- Just-in-time compilers (e.g., Sun JVM, IBM JikesVM)
 - bytecode is compiled to native code at load time (or later)
 - code can be optimized (at compile time or later)
 - quicker
 - slow startup
- Ahead-of time compilers (e.g., GCJ)
 - bytecode is compiled to native code offline
 - quick startup
 - quick execution
 - static compilation

JVM: Loader – The Classfile Format

```
ClassFile {  
    version  
    constant pool  
    flags  
    super class  
    interfaces  
    fields  
    methods  
    attributes  
}
```

Constants:

- Values
String / Integer / Float / ...
- References
Field / Method / Class / ...

Attributes:

- ConstantValue
- Code
- Exceptions

JVM: Class File Format

```
class HelloWorld {  
  
    public static void printHello() {  
        System.out.println("hello, world");  
    }  
  
    public static void main (String[] args) {  
        HelloWorld myHello = new HelloWorld();  
        myHello.printHello();  
    }  
  
}
```

JVM: Class File (Constant Pool)

- | | | | |
|-----------------|--------------------------------|-------------|------------------------|
| 1. String | hello, world | 15. Unicode | ()V |
| 2. Class | HelloWorld | 16. Unicode | (Ljava/lang/String;)V |
| 3. Class | java/io/PrintStream | 17. Unicode | |
| 4. Class | java/lang/Object | | ([Ljava/lang/String;)V |
| 5. Class | java/lang/System | 18. Unicode | <init> |
| 6. Methodref | HelloWorld.<init>() | 19. Unicode | Code |
| 7. Methodref | | 20. Unicode | ConstantValue |
| | java/lang/Object.<init>() | 21. Unicode | Exceptions |
| 8. Fieldref | java/io/PrintStream | 22. Unicode | HelloWorld |
| | java/lang/System.out | 23. Unicode | HelloWorld.java |
| 9. Methodref | | 24. Unicode | LineNumberTable |
| | HelloWorld.printHello() | 25. Unicode | Ljava/io/PrintStream; |
| 10. Methodref | | 26. Unicode | LocalVariables |
| | java/io/PrintStream.println(ja | 27. Unicode | SourceFile |
| | va/lang/String) | 28. Unicode | hello, world |
| 11. NameAndType | <init> ()V | 29. Unicode | java/io/PrintStream |
| 12. NameAndType | out | 30. Unicode | java/lang/Object |
| | Ljava/io/PrintStream; | 31. Unicode | java/lang/System |
| 13. NameAndType | printHello ()V | 32. Unicode | main |
| 14. NameAndType | println | 33. Unicode | out |
| | (Ljava/lang/String;)V | 34. Unicode | printHello |

JVM: Class File (Code)

Methods

```
0 <init>()
  0 ALOAD0
  1 INVOKESPECIAL [7]  java/lang/Object.<init>()
  4 RETURN

1 PUBLIC STATIC main(java/lang/String [])
  0 NEW [2]  HelloWorld
  3 DUP
  4 INVOKESPECIAL [6]  HelloWorld.<init>()
  7 ASTORE1
  8 INVOKESTATIC [9]  HelloWorld.printHello()
 11 RETURN

2 PUBLIC STATIC printHello()
  0 GETSTATIC [8]  java/io/PrintStream java/lang/System.out
  3 LDC1  hello, world
  5 INVOKEVIRTUAL [10]  java/io/PrintStream.println(java/lang/String )
  8 RETURN
```

JVM: Compilation – Pattern Expansion

- Each byte code is translated according to fix patterns
 - + easy
 - limited knowledge
- Example (pseudocode)

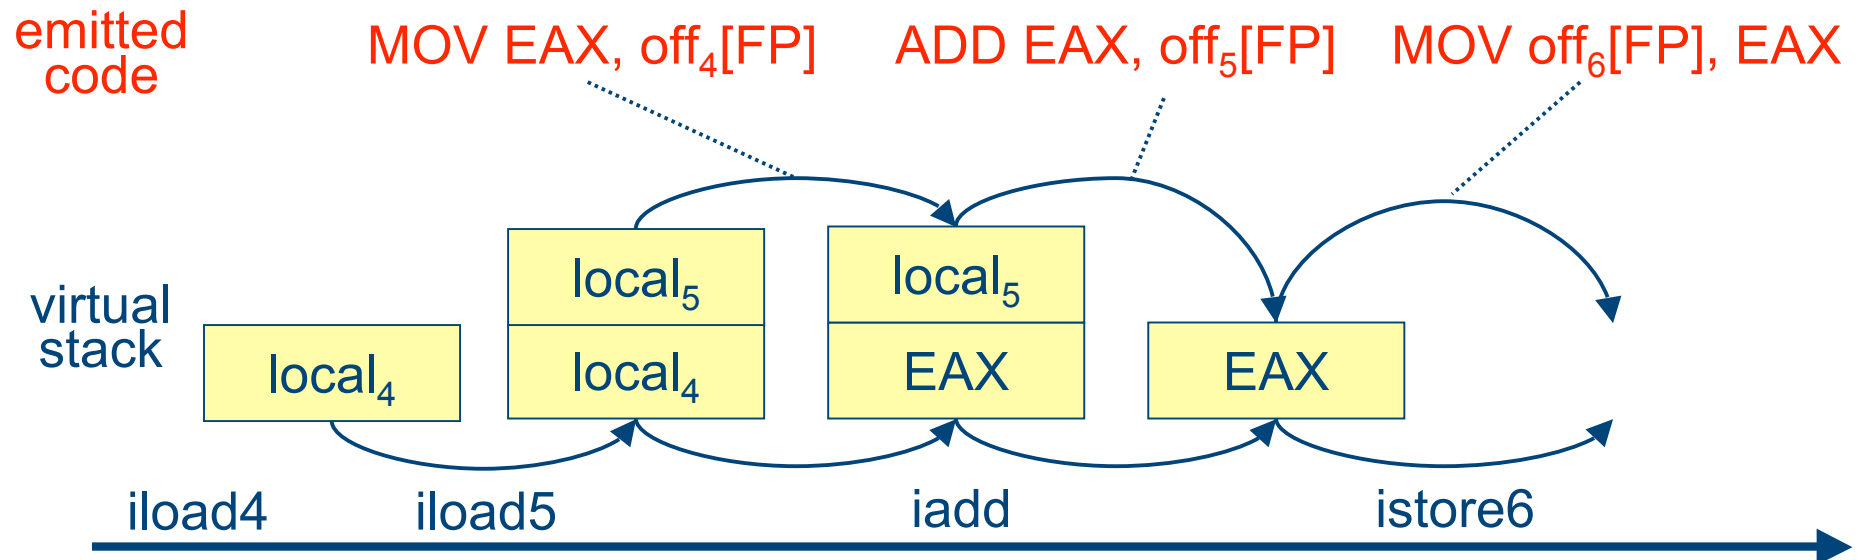
```
switch (o) {  
case ICONST<n>: generate("push n"); PC++; break;  
case ILOAD<n>: generate("push off_n[FP]"); PC++; break;  
case IADD: generate("pop -> R1");  
            generate("pop -> R2");  
            generate("add R1, R2 -> R1");  
            generate("push R1");  
            PC++;  
            break;  
...  
}
```

JVM: Optimizing Pattern Expansion

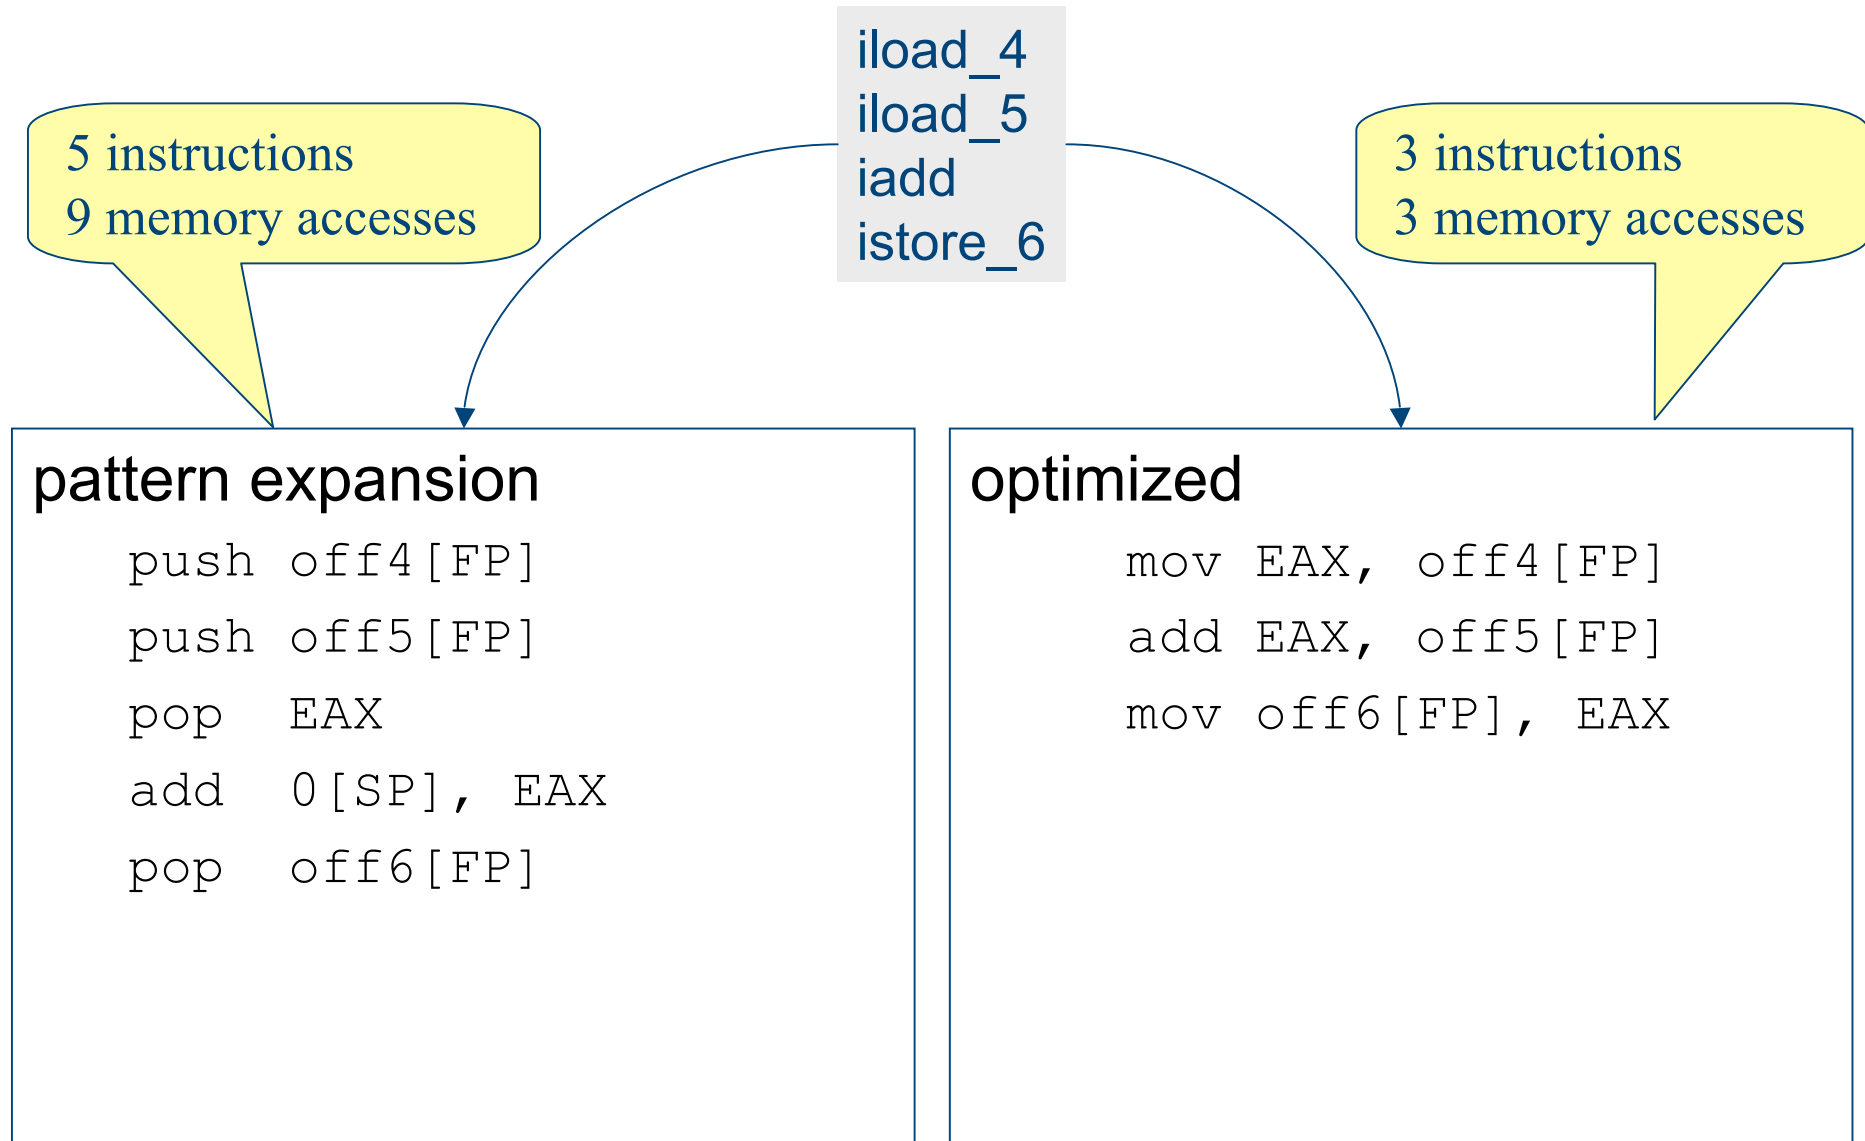
Main Idea:

- use internal **virtual stack**
- stack values are consts / fields / locals / array fields / registers / ...
- flush stack **as late as possible**

```
iload 4  
iload 5  
iadd  
istore 6
```




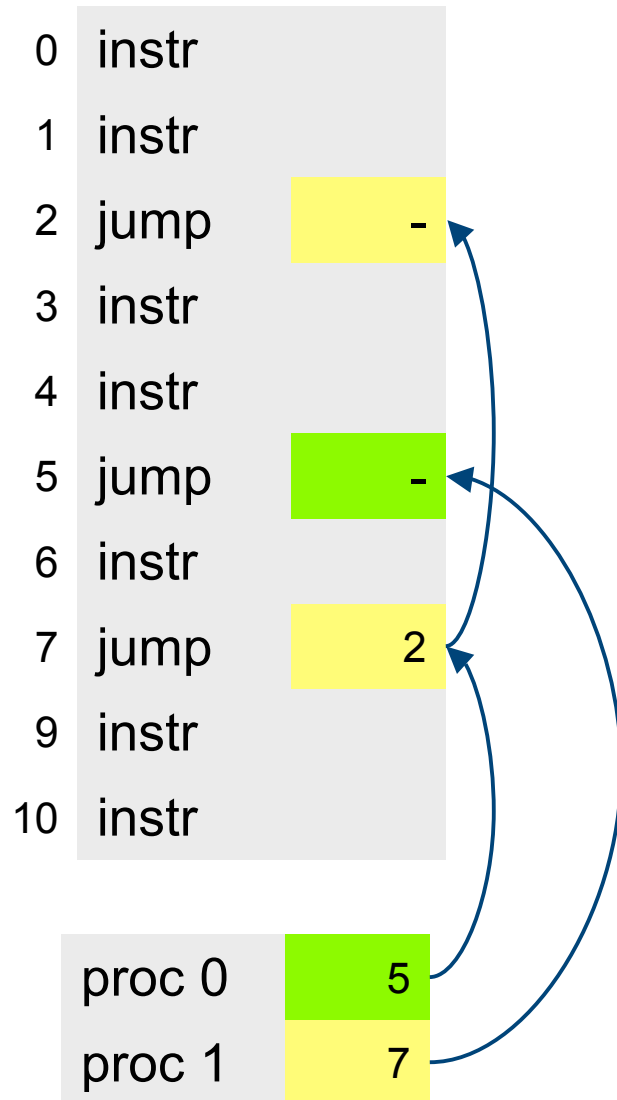
JVM: Compiler Comparison



Linking (General)

- A compiled program contains references to **external code** (libraries)
- After loading the code the system need to **link** the code to the library
 - identify the calls to external code
 - locate the callees (and load them if necessary)
 - patch the loaded code
- Two options:
 - the code contains a list of sites for each callee
 - the calls to external code are jumps to a *procedure linkage table* which is then patched (double indirection)


Linking (General)



Linking (General)

0	instr	
1	instr	
2	jump	&p1
3	instr	
4	instr	
5	jump	&p0
6	instr	
7	jump	&p1
9	instr	
10	instr	

proc 0	5
proc 1	7



JVM: Linking



- Bytecode interpreter
 - references to other objects are made through the JVM (e.g., invokevirtual, getfield, ...)
- Native code (ahead of time compiler)
 - static linking
 - classic native linking
- JIT compiler
 - only some classes are compiled
 - calls could reference classes that are not yet loaded or compiled (delayed compilation)
 - ➔ code instrumentation

JVM: Methods and Fields Resolution


- method and fields are accessed through special VM functions (e.g., `invokevirtual`, `getfield`, ...)
- the parameters of the special call defines the target
- the parameters are indexes in the constant pool
- the VM checks if the call is legal and if the target is present

JVM: JIT – Linking and Instrumentation


- Use code instrumentation to detect first access of static fields and methods



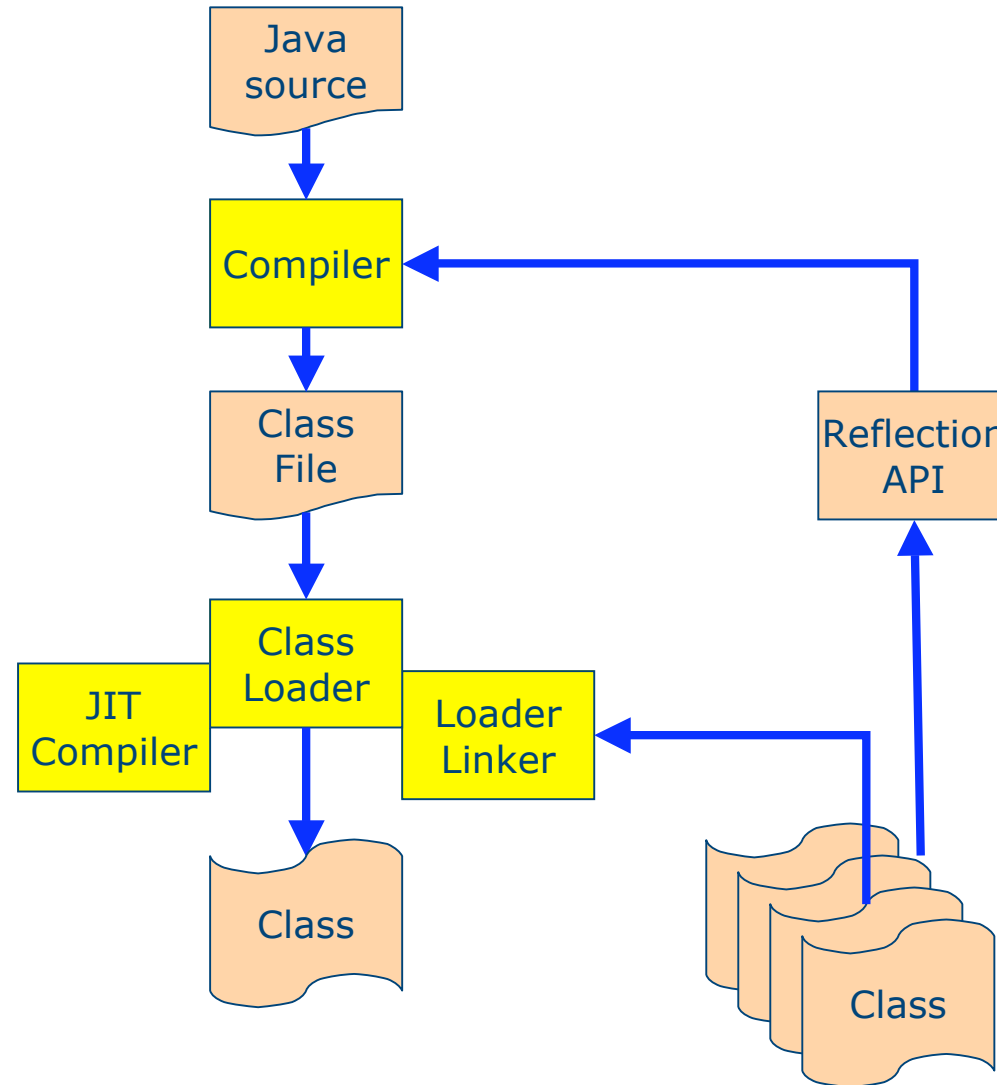
Compilation and Linking Overview



Compilation and Linking Overview



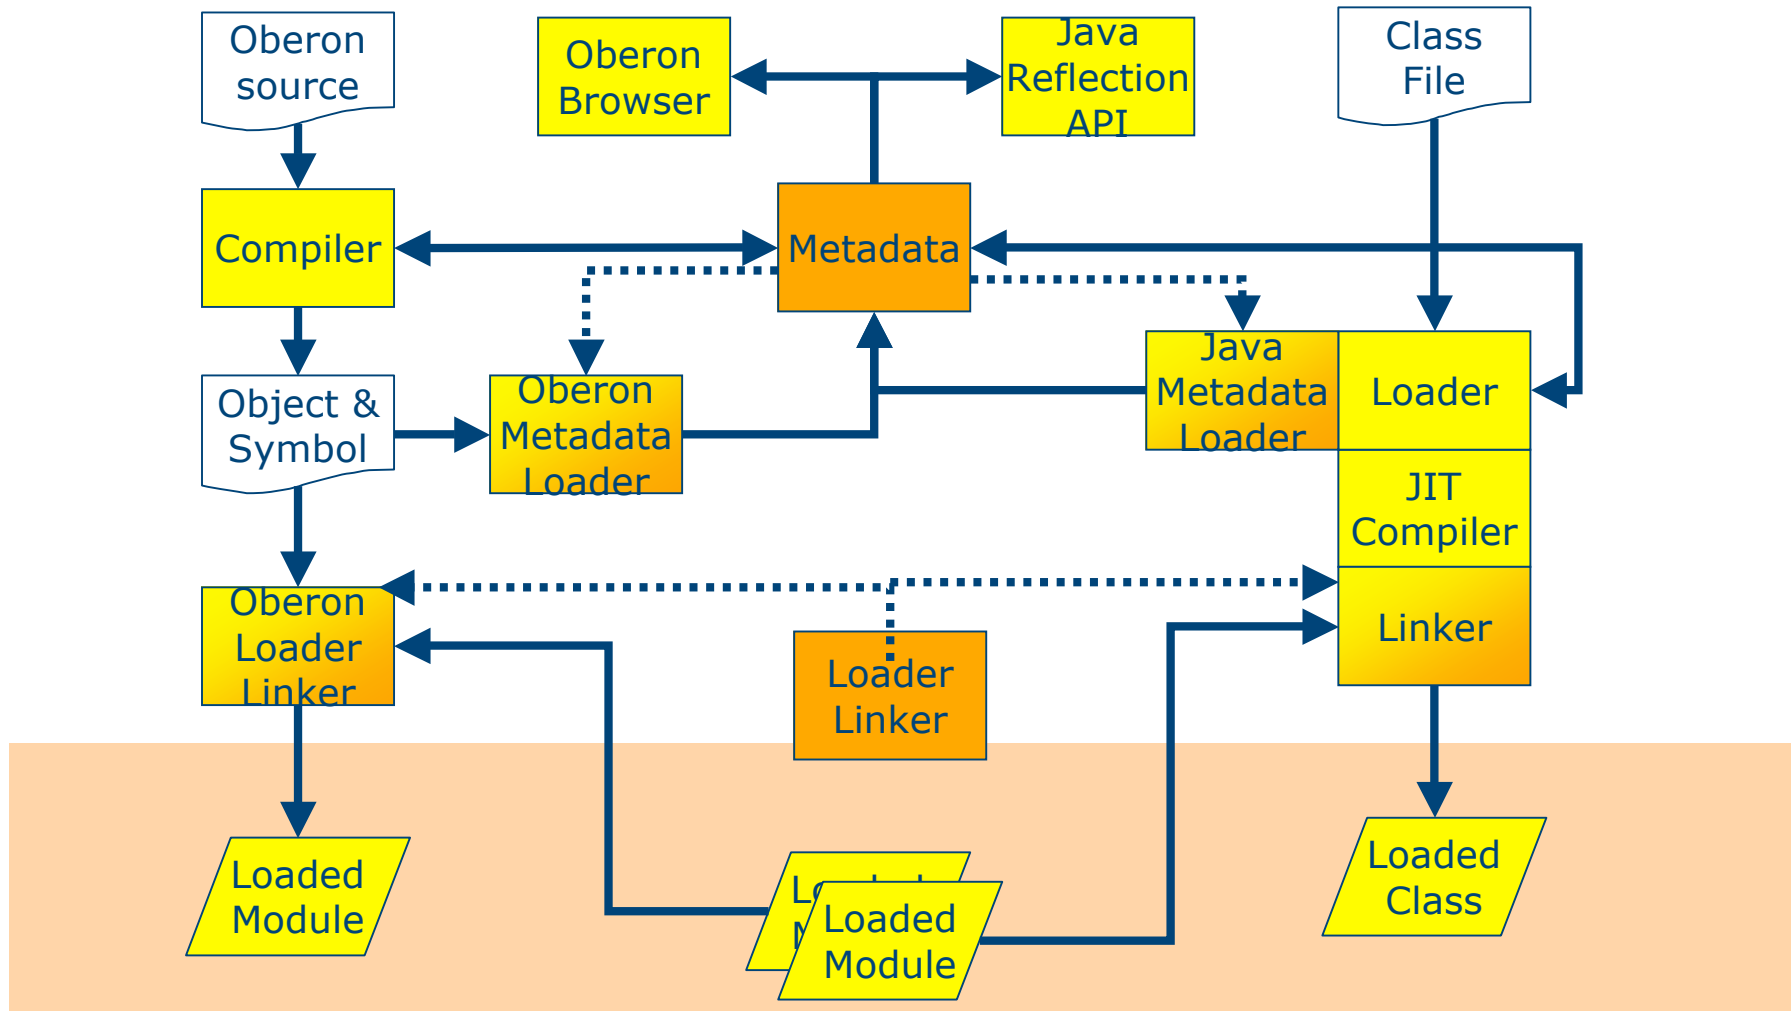
Compilation and Linking Overview



Jaos

- Jaos (**J**ava on **A**ctive **O**bject **S**ystem) is a Java virtual machine for the Bluebottle system
- goals:
 - implement a JVM for the Bluebottle system
 - show that the Bluebottle kernel is generic enough to support more than one system
 - interoperability between the Active Oberon and Java languages
 - interoperability between the Oberon System and the Java APIs

Jaos (Interoperability Framework)



JVM: Verification

- Compiler generates “good” code....
 - that could be changed before reaching the JVM
- ➔ need for verification

Verification makes the VM simpler (less run-time checks):

- no operand stack overflow
- load / stores are valid
- VM types are correct
- no pointer forging
- no violation of access restrictions
- access objects as they are (type)
- local variable initialized before load
- ...

JVM: Verification

Pass1 (Loading):

- class file version check
- class file format check
- class file complete

Pass 2 (Linking):

- final classes are not subclassed
- every class has a superclass (but Object)
- constant pool references
- constant pool names

JVM: Verification

Delayed for
performance
reasons

Pass 3 (Linking):

For each operation in code
(independent of the path):

- operation stack size is the same
- accessed variable types are correct
- method parameters are appropriate
- field assignment with correct types
- opcode arguments are appropriate

Byte-Code
Verification

Pass 4 (RunTime):

First time a type is referenced:

- load types when referenced
- check access visibility
- class initialization

First member access:

- member exists
- member type same as declared
- current method has right to access member

JVM: Byte-Code Verification

Verification:

- branch destination must exist
- opcodes must be legal
- access only existing locals
- code does not end in the middle of an instruction
- types in byte-code must be respected
- execution cannot fall off the end of the code
- exception handler begin and end are sound

JVM: Bootstrapping

How to start a JVM?

- External help needed!
- Load core classes
- Compile classes
- Provide memory management
- Provide threads

Solution:

Implement Java on 3rd party system

- Linux
- Solaris
- Windows
- Bluebottle
- Java

Bootstrapping: Jaos Example

- All native methods in Active Oberon
- Use Bluebottle run-time structures
 - module (= class)
 - type descriptor
 - object (= object instance)
 - active object (= thread)
- Bootstrap
 - load core classes
 - Object, String, System, Runtime, Threads, ...
 - Exception
 - forward exception to java code
 - allocate java classes from Oberon

Bootstrapping: Jnode VM Example

- JVM written in Java
- small core in assembler
 - low-level functionalities that requires special assembler instructions
- some native methods inlined by the compiler
 - Unsafe
 - debug(String)
 - int AddressToInt(Address)
 - int getInt(Object, offset)
- Bootstrap
 - compile to Java classes
 - bootloader:
 - native compilation
 - code placement
 - structure allocation
 - make boot image
 - boot with GNU/GRUB

Bootstrapping: Oberon / Bluebottle

- Compile each module to machine code
 - system calls for
 - newrec (record)
 - newsys (block, no ptrs)
 - newarr (array)
 - linker / bootlinker patches syscalls with real procedures
- bootlinker is same as linker, but uses different memory addresses
 - simulates memory allocation
 - start address configurable
 - glue code to call all module bodies

Bootstrapping Compilers

